

# Small and Easy Motivation

```
#include <stdio.h>
#include <stdlib.h>

int main() {

    // Do this part in parallel

    printf( "Hello, World!\n" );

    return 0;
}
```

# Small and Easy Motivation

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main() {

    omp_set_num_threads(16);

    // Do this part in parallel
    #pragma omp parallel
    {
        printf( "Hello, World!\n" );
    }

    return 0;
}
```

# Simple!

Serial Program:

```
void main()
{
    double Res[1000];

    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

Parallel Program:

```
void main()
{
    double Res[1000];
    #pragma omp parallel for
    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

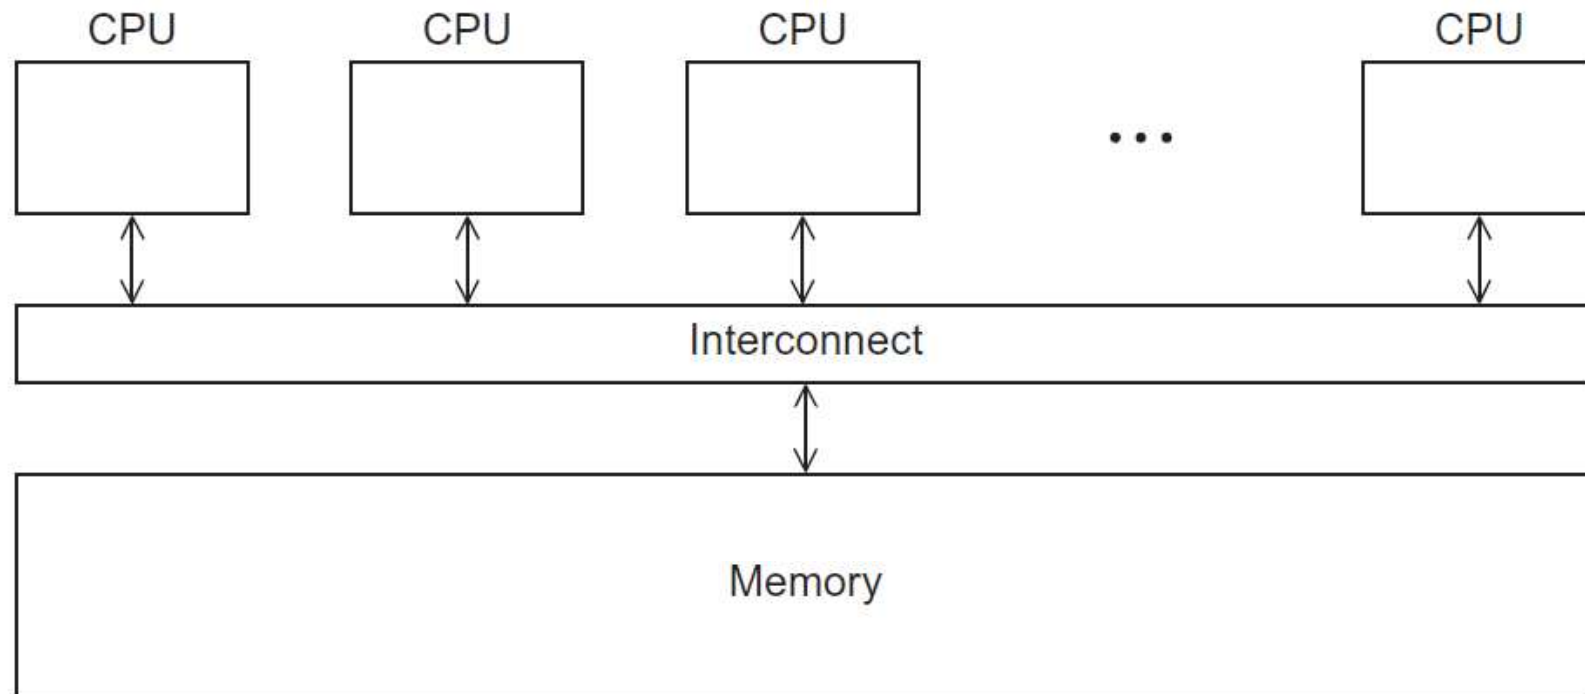
OpenMP can parallelize many serial programs with relatively few annotations that specify parallelism and independence

OpenMP is a small API that hides cumbersome threading calls with simpler *directives*

# OpenMP

- An **API** for shared-memory parallel programming.
- Designed for systems in which each **thread** can potentially have access to all available memory.
- System is viewed as a collection of cores or CPU's, all of which have access to the same main memory → shared memory architecture

# A shared memory system



# Pragmas

- Special **preprocessor** instructions.
- specified by the C standard for providing additional information to the compiler, beyond what is conveyed in the language itself.
- Compilers that don't support the pragmas ignore them.

**#pragma**

```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void);  /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);

    # pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
}  /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);

}  /* Hello */




```

```
gcc -g -Wall -fopenmp -o omp_hello omp_hello .c
```

```
./omp_hello 4
```

 running with 4 threads

Hello from thread 0 of 4  
Hello from thread 1 of 4  
Hello from thread 2 of 4  
Hello from thread 3 of 4

 possible  
outcomes   


Hello from thread 1 of 4  
Hello from thread 2 of 4  
Hello from thread 0 of 4  
Hello from thread 3 of 4


Hello from thread 3 of 4  
Hello from thread 1 of 4  
Hello from thread 2 of 4  
Hello from thread 0 of 4



# OpenMp pragmas

- `# pragma omp parallel`
  - Most basic parallel directive.
  - The number of threads that run the following `structured block of code` is determined by the run-time system.

# A process forking and joining two threads



# clause

- Text that modifies a directive.
- The num\_threads clause can be added to a parallel directive.
- It allows the programmer to specify the number of threads that should execute the following block.


```
# pragma omp parallel num_threads ( thread_count )
```

## Of note...


- There may be system-defined limitations on the number of threads that a program can start.
- The OpenMP standard doesn't guarantee that this will actually start `thread_count` threads.
- Unless we're trying to start a lot of threads, we will almost always get the desired number of threads.

# Some terminology

- In OpenMP parlance the collection of threads executing the parallel block — the original thread and the new threads — is called a **team**, the original thread is called the **master**, and the additional threads are called **slaves**.



# Again: The trapezoidal rule



# Serial algorithm

```
/* Input:  a, b, n */  
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++) {  
    x_i = a + i*h;  
    approx += f(x_i);  
}  
approx = h*approx;
```

# A First OpenMP Version


- 1) We identified two types of tasks:
  - a) computation of the areas of individual trapezoids, and
  - b) adding the areas of trapezoids.
- 2) There is no communication among the tasks in the first collection, but each task in the first collection communicates with task 1b.



# A First OpenMP Version

- 3) We assumed that there would be many more trapezoids than cores.
- So we aggregated tasks by assigning a contiguous block of trapezoids to each thread.

# Assignment of trapezoids to threads



Time	Thread 0	Thread 1
0	<code>global_result = 0 to register</code>	<code>finish my_result</code>
1	<code>my_result = 1 to register</code>	<code>global_result = 0 to register</code>
2	<code>add my_result to global_result</code>	<code>my_result = 2 to register</code>
3	<code>store global_result = 1</code>	<code>add my_result to global_result</code>
4		<code>store global_result = 2</code>

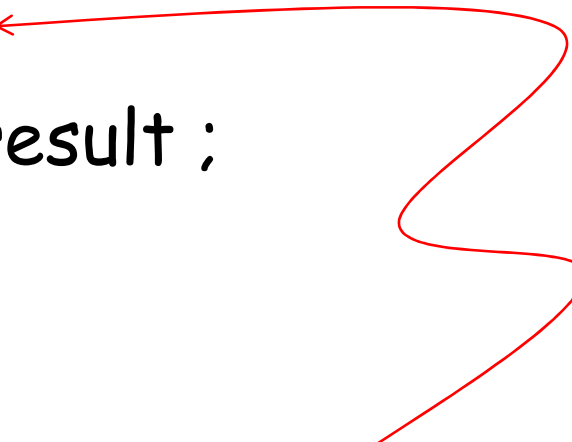
Unpredictable results when two (or more) threads attempt to simultaneously execute:

`global_result += my_result ;`



# Mutual exclusion

```
# pragma omp critical  
global_result += my_result ;
```



only one thread can execute  
the following structured block at  
a time

```

#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Trap(double a, double b, int n, double* global_result_p);

int main(int argc, char* argv[]) {
    double    global_result = 0.0;  /* Store result in global_result */
    double    a, b;                 /* Left and right endpoints      */
    int       n;                    /* Total number of trapezoids    */
    int       thread_count;

    thread_count = strtol(argv[1], NULL, 10);
    printf("Enter a, b, and n\n");
    scanf("%lf %lf %d", &a, &b, &n);
    # pragma omp parallel num_threads(thread_count)
    Trap(a, b, n, &global_result);

    printf("With n = %d trapezoids, our estimate\n", n);
    printf("of the integral from %f to %f = %.14e\n",
        a, b, global_result);
    return 0;
} /* main */

```

```

void Trap(double a, double b, int n, double* global_result_p) {
    double h, x, my_result;
    double local_a, local_b;
    int i, local_n;
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    h = (b-a)/n;
    local_n = n/thread_count;
    local_a = a + my_rank*local_n*h;
    local_b = local_a + local_n*h;
    my_result = (f(local_a) + f(local_b))/2.0;
    for (i = 1; i <= local_n-1; i++) {
        x = local_a + i*h;
        my_result += f(x);
    }
    my_result = my_result*h;

    # pragma omp critical
    *global_result_p += my_result;
} /* Trap */

```

# Scope

- In serial programming, the scope of a variable consists of those parts of a program in which the variable can be used.
- In OpenMP, the scope of a variable refers to the set of threads that can access the variable in a parallel block.

# Scope in OpenMP

- A variable that can be accessed by all the threads in the team has **shared** scope.
- A variable that can only be accessed by a single thread has **private** scope.
- The default scope for variables declared before a parallel block is shared.





```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

void Hello(void); /* Thread function */

int main(int argc, char* argv[]) {
    /* Get number of threads from command line */
    int thread_count = strtol(argv[1], NULL, 10);

    # pragma omp parallel num_threads(thread_count)
    Hello();

    return 0;
} /* main */

void Hello(void) {
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    printf("Hello from thread %d of %d\n", my_rank, thread_count);
} /* Hello */
```

shared

private

```

void Trap(double a, double b, int n, double* global_result_p) {
    double h, x, my_result;
    double local_a, local_b;
    int i, local_n;
    int my_rank = omp_get_thread_num();
    int thread_count = omp_get_num_threads();

    h = (b-a)/n;
    local_n = n/thread_count;
    local_a = a + my_rank*local_n*h;
    local_b = local_a + local_n*h;
    my_result = (f(local_a) + f(local_b))/2.0;
    for (i = 1; i <= local_n-1; i++) {
        x = local_a + i*h;
        my_result += f(x);
    }
    my_result = my_result*h;

    # pragma omp critical
    *global_result_p += my_result;
}.../*..Trap. */

```

Do you remember  
the trapezoidal?

We need this more complex version to add each thread's local calculation to get *global\_result*.

```
void Trap(double a, double b, int n, double* global_result_p);
```

Although we'd prefer this.

```
double Trap(double a, double b, int n);
```



```
global_result = Trap(a, b, n);
```

How about this:

```
double Local_trap(double a, double b, int n);
```

and we use it like this:

```
global_result = 0.0;  
# pragma omp parallel num_threads(thread_count)  
{  
#   pragma omp critical  
    global_result += Local_trap(double a, double b, int n);  
}
```

... we force the threads to execute sequentially.

It is now slower than a version with single thread!

How can we fix this?

We can avoid this problem by:

1. declaring a private variable inside the parallel block
2. moving the critical section after the function call

```
global_result = 0.0;
# pragma omp parallel num_threads(thread_count)
{
    double my_result = 0.0;  /* private */

    my_result += Local_trap(double a, double b, int n);
# pragma omp critical
    global_result += my_result;
}
```

Can we do better?

# Reduction operators

- Is a binary operation
- A **reduction** is a computation that repeatedly applies the same reduction operator to a sequence of operands in order to get a single result (we have seen that in MPI!).
- All of the intermediate results of the operation should be stored in the same variable: the **reduction variable**.

A reduction clause can be added to a parallel directive.

```
reduction(<operator>: <variable list>)
```



`+, *, -, &, |, ^, &&, ||`

Be careful of:

- subtraction
- floating points

And the code becomes:

```
global_result = 0.0;  
# pragma omp parallel num_threads(thread_count) \  
  reduction(+: global_result)  
  global_result += Local_trap(double a, double b, int n);
```

# How Does OpenMP Do it?


- The reduction variable is shared.
- OpenMP create a local variable for each thread
- Those local variables are initialized to the identity value for the reduction operator
- When the parallel block ends, the values in the private variables are combined into the shared variable.



# #pragma omp parallel for

- Forks a team of threads to execute the following structured block.
- The structured block following the parallel for directive **must be a for loop**.
- The system parallelizes the for loop by **dividing the iterations of the loop among the threads**.

```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
for (i = 1; i <= n-1; i++)  
    approx += f(a + i*h);  
approx = h*approx;
```



```
h = (b-a)/n;  
approx = (f(a) + f(b))/2.0;  
# pragma omp parallel for num_threads(thread_count) \  
    reduction(+: approx)  
    for (i = 1; i <= n-1; i++)  
        approx += f(a + i*h);  
approx = h*approx;
```

In a loop that is parallelized with *parallel for* the default scope of a loop variable is private

# Legal forms for parallelizable *for statements*

<b>for</b>	{	index = start ;		index++
				++index
			index < end	index--
			index <= end	--index
			index >= end ;	index += incr
			index > end	index -= incr
				index = index + incr
				index = incr + index
		index = index - incr	)	

Number of iterations **MUST** be known  
prior to the loop execution.

There can be a call to *exit* in the loop body.

OpenMP won't parallelize while loops or do-while loops.

# Caveats

- The variable **index** must have integer or pointer type (e.g., it can't be a float).
- The expressions **start**, **end**, and **incr** must have a compatible type. For example, if **index** is a pointer, then **incr** must have integer type.


# Caveats

- The expressions **start**, **end**, and **incr** must not change during execution of the loop.
- During execution of the loop, the variable **index** can only be modified by the "increment expression" in the for statement.

# Data dependencies

```
fibonacci[0] = fibonacci[1] = 1;  
for (i = 2; i < n; i++)  
    fibonacci[i] = fibonacci[i-1] + fibonacci[i-2];
```

note 2 threads



```
fibonacci[0] = fibonacci[1] = 1;  
# pragma omp parallel for num_threads(2)  
for (i = 2; i < n; i++)  
    fibonacci[i] = fibonacci[i-1] + fibonacci[i-2];
```

1 1 2 3 5 8 13 21 34 55

this is correct

but sometimes  
we get this

1 1 2 3 5 8 0 0 0 0

# What happened?

1. OpenMP compilers don't check for dependences among iterations in a loop that's being parallelized with a parallel for directive.
2. A loop in which the results of one or more iterations depend on other iterations cannot, in general, be correctly parallelized by OpenMP.



# Question

Do we have to worry about the following:

```
#pragma omp parallel for num_threads(2)
  for( i = 0 ; i < n; i++) {
    x[i] = a + i*h;
    y[i] = exp(x[i]);
  }
```